

Task Discovery Service (TDS): A Dual-Architecture Service Registry for Distributed Systems

Final Year Project Report

February 2026

1 Introduction and Motivation

In large distributed systems, one weak point can force the whole platform to slow down: static service routing. The core issue in the system I worked on at Cubic Transportation Systems was not input handling, but rigid, hard-coded output endpoints. Although the component accepted requests from a dynamic number of devices, every routing change required redevelopment and full regression testing. This made scaling expensive and slow, and over time the original resolvers had to be replaced with load balancers as demand outgrew the fixed design. This project began from that operational bottleneck and asks how service discovery can be made both dynamic and scalable.

1.1 Problem Statement

The problem lay in two key aspects

1. The component was not dynamic even though it performed the same function for any given request from the devices beneath it.
2. The server was mainly overloaded because it accepted the data from the requesting device and passed it on to the resolver. The data was never operated upon or even observed by the component.

I proposed to address these issues by designing an architecture that would perform and scale with the system.

1.2 Research Objectives

This project addresses the service discovery problem by developing the Task Discovery Service (TDS), a flexible registry system that supports both centralised and peer-to-peer (P2P) operational modes within a unified architecture. The primary objectives are:

1. Design and implement a dual-mode service registry supporting both centralised and P2P architectures
2. Develop efficient algorithms for service registration, query routing, and load balancing
3. Create a comprehensive monitoring interface for real-time system visualisation
4. Evaluate performance characteristics under various network conditions and load patterns
5. Provide multiple transport and security options including UDP, TCP, and mutual TLS authentication

1.3 Contribution

The main contribution of this work is a production-ready service discovery system that eliminates the operational mode dichotomy present in existing solutions. TDS allows development teams to begin with a simple centralised deployment and seamlessly transition to distributed P2P operation as scale requirements increase, without changing client code or protocol definitions. This architectural flexibility, combined with comprehensive security features and real-time monitoring capabilities, represents a significant advancement in practical service discovery systems.

2 Project Background and Context

2.1 Service Discovery in Distributed Systems

Service discovery allows one part of a distributed system to find another at runtime. Instead of hard-coding IP addresses or ports, services register themselves and clients query for the address they need. This matters because modern systems change constantly: instances are restarted, scaled up, moved, or fail. If the discovery mechanism is out of date, the rest of the system may fail even when the application logic is correct.

Most existing solutions fall into three broad groups: centralised registries, decentralised peer-to-peer systems, and hybrid systems that combine a registry with programmable traffic control. All three solve the same basic problem, each having their different advantages and drawbacks.

2.1.1 Centralised registry approaches

Centralised approaches keep a single authoritative view of the system. Services register with a central catalogue and clients query that catalogue through an API or DNS interface. Systems such as etcd and other Raft-based control planes are designed to give fresh, strongly consistent control-plane state [3, 13, 20]. This makes them easier to reason about because there is one place where the truth lives.

Many real systems use this model. Kubernetes keeps a stable service identity while its EndpointSlices track changing pods, and cluster DNS translates the service name to current backends [16]. Consul follows a similar pattern, but adds health checks and multiple discovery methods [2, 19].

The main strength of centralised discovery is clarity. It is usually easier to manage, observe, and secure. The main weakness is dependency on the control plane itself. If the registry becomes overloaded or unavailable, discovery becomes a bottleneck.

2.1.2 Decentralised approaches

Decentralised approaches remove the single authority and spread responsibility across many peers. In this family, lookup data is distributed around the network rather than stored in one registry. Systems such as Chord and Kademlia showed that this can scale well by placing data according to a hash of the key, in this case a service name, and routing queries through the network efficiently [1, 10]. Dynamo extended this style of thinking by favouring availability and replication even when the network is unstable [11].

These systems also need a way to detect which peers are alive. Protocols such as SWIM address this through scalable membership and failure detection [12]. At smaller scale, mDNS and DNS-SD provide a simpler decentralised model for local networks, where devices can advertise and discover services without a dedicated server [14, 15].

The strength of decentralised discovery is that there is no single point of failure. The weakness is that it is harder to guarantee that every node has the same up-to-date view of the system. As the network grows, correctness depends on how well peer information spreads.

2.1.3 Traffic-management, load-balancing, and policy-enforcement approaches

A third family focuses less on the problem of *finding* services and more on the problem of *choosing* between available instances and *controlling* which connections are allowed. These approaches usually sit on top of an existing discovery mechanism and add traffic-management or policy-enforcement behaviour.

For load balancing, systems such as Consul and Istio use discovery information to route traffic only to healthy instances and to apply routing strategies beyond simple name resolution [18, 19]. Consul can expose healthy endpoints through DNS and integrate with external load balancers, while Istio applies policies such as least-requests, round-robin, weighted routing, retries, timeouts, and circuit breaking through proxies. This makes them useful when the main challenge is not just locating a service, but distributing load safely and efficiently across many replicas.

For network control, platforms such as Kubernetes provide policy layers that restrict which workloads may communicate at all [17]. NetworkPolicy rules can isolate ingress and egress traffic and allow only selected pods, namespaces, ports, or IP ranges. In effect, these systems act like structured, application-aware firewall mechanisms inside a distributed platform.

The advantage of this family is control. It gives operators a way to combine service discovery with health-aware routing, traffic shaping, and access restrictions. The disadvantage is that it depends on additional infrastructure and configuration. These solutions are powerful, but they can become difficult to reason about because discovery, load balancing, and policy enforcement are spread across several layers rather than handled in one place.

2.2 Security, Reachability, and Policy Correctness

Even with correct discovery and load balancing, communication can still fail if policy blocks the connection. This is why service discovery must be considered alongside reachability and access control.

The firewall literature is useful here. Firewall Policy Queries focuses on detecting rule anomalies such as shadowing and redundancy, while Model Checking Firewall Policy Configurations focuses on verifying that policy behaves as intended [21, 22]. These works reinforce a simple point: a discovery system is only fully useful when the endpoints it returns are not just correct, but also reachable under the active security policy.

2.3 How TDS Differs, and Its Strengths and Weaknesses

Existing solutions usually require an early commitment to one style of architecture. A team must choose a centralised registry, a decentralised overlay, or a more complex hybrid platform and then build around that choice. TDS differs by supporting both centralised and peer-to-peer operation behind the same client-facing interface.

The main design idea is that services do not interact directly with the registry implementation. Instead, they talk to a local client proxy using the same REGISTER and QUERY operations in both modes. This means the deployment architecture can change without forcing application code to change with it. In practical terms, a system can begin with centralised discovery, which is simpler to operate and easier to observe, and later move to P2P operation if resilience or removal of a central dependency becomes more important.

This gives TDS several inherent strengths. First, it reduces lock-in to a single discovery architecture. Second, it lowers adoption cost because the service-facing interface remains small and stable. Third, it allows the centralised mode to benefit from straightforward control and governance, while the P2P mode can offer better tolerance of single-node failure. Finally, the design leaves room for policy-aware routing, which is important in environments where not every discovered service should be reachable by every requester.

However, this flexibility also introduces weaknesses. Supporting two modes increases implementation complexity because both architectural paths must be maintained and tested. The centralised mode still inherits the usual problem of depending on a core registry service. The P2P mode avoids that single dependency, but it is naturally weaker in terms of global consistency and can become harder to reason

about as peer knowledge diverges. In other words, TDS avoids early architectural lock-in, but it does not remove the underlying trade-offs between centralised and decentralised discovery; it exposes both and gives the user a choice.

For that reason, TDS is best understood not as a replacement for every mature discovery platform, but as a practical bridge between architectural extremes. Its contribution is to make movement between those extremes possible without changing the service-side contract.

3 Design

3.1 services, tasks, and the client proxy

The diagram below gives an overview of a client machine in this system. With multiple services registering and querying services to the client proxy which is the local endpoint of this solution.




Figure 1: Client machine service registration and discovery via client proxy

In this architecture, services running on a machine will query their local instance of the client proxy. As far as the services are concerned, that is all that needs to be done for them to advertise their service along with sending periodic heartbeats. The client proxy will also handle any requests these services have for where to send their data once they have finished the scope of their operations on it. The client proxy is the interface behind which all the mechanisms of this project lie.

3.1.1 Message format

All communication between local services (clients) and the client proxy uses JSON messages in both centralised and P2P modes. The two shared commands are **REGISTER** and **QUERY**.

JSON requests (service → client proxy)

```
{ "cmd": "REGISTER", "task": "payment", "address": "10.0.1.5:8080", "capacity": 10 }
{ "cmd": "QUERY", "task": "payment" }
```

JSON response format

```
{ "status": "OK" } // Register - OK
{ "status": "OK", "address": "10.0.1.5:8080" } // Successful query
{ "status": "NOTFOUND" } // queried task not found
{ "status": "ERR", "error": "task required" } // example error: no task in query request
```

These message types are intentionally minimal so the same client-side behaviour works regardless of whether the proxy is currently forwarding to a central server or to the P2P registry.

3.2 Centralised mode

We now move outside the scope of the single machine registering and querying services to a distributed system containing many such machines. The below figure demonstrates the role the server plays in centralised mode.




Figure 2: Client machines operating in server architecture

3.2.1 load balancing

The server load balances services using a round-robin mechanism for all services under a task. This ensures even distribution of traffic to each service. Configurable on top of this is a weighted round robin mechanism using each registered service's capacity as a weighting. This ensures that services running on less powerful machines are not swamped by being routed the same volume of traffic as a more powerful variant. During the heartbeat message, a service advertises its current capacity to inform load balancing.

3.2.2 Hybrid storage: cache and persistence

The centralised server uses a hybrid storage strategy that balances performance with operational durability. At the core is an in-memory cache that holds a subset of frequently queried tasks. Upon a cache miss or at regularised intervals, the server warms the cache by querying a persistent PostgreSQL backend. This approach achieves low-latency query response for hot tasks while maintaining durable, databaseable state for recovery and operational continuity.

In practice, the server maintains a configurable cache size and uses query frequency to determine which tasks remain cached. After query count thresholds or time-based intervals, the cache is refreshed by querying the database for the current service entries. This ensures that infrequently accessed tasks do not consume valuable in-memory space, and that heartbeat updates in the database are reflected in the next cache warm before they are served to clients. The hybrid approach proved essential to achieving good query latency without sacrificing the ability to scale the deployment by adding a database backend and persisting across restart events.

3.2.3 Firewall-mode

In many real use cases of this project, while all machines are able to query and register with the TDS, not all services are allowed to communicate with each other. This is enforced by networking rules that are created and managed by a central authority in the system. The purpose of firewall-mode is to cross-reference a query for a service with the list of ip-address:ports that the requestor can communicate with. This way the server will only return a service address that the client can use.

3.3 Peer-to-Peer mode

The diagram below gives an overview of the peer to peer architecture. It shows how tasks are registered and then stored on the correct node in the network. It also demonstrates service queries and then the separate sending of data outside of the architecture for the process. In contrast to the centralised mode,




Figure 3: Client machines operating in P2P architecture

all functionality of this architecture runs from the client-proxy. The P2P network is formed from one "bootstrap" node. All subsequent nodes that join are given a pre-existing node in the network as a bootstrap. When a node joins, it hashes its ip address to receive their node ID. This ID is capped at 256 bits forming a ring network.

3.3.1 Distributed Hash Table (DHT)

The DHT is used to store the task:address mappings in this architecture. When a service registers a task with their client-proxy, the proxy hashes the service name. It then searches its list of peers to find the node ID that is closest to this hash. Once that node is found the client-proxy sends a STORE request to it with the service details. The node that receives this store message compares its peer list to the task hash to ensure it is the most appropriate candidate. If so then it stores the task in memory similarly to the cache in centralised mode.

A query follows the same path: the client-proxy hashes the query task, sends a query request to the node with the closest ID. If that node has that service stored it returns an address, if not then it returns NOTFOUND.

3.3.2 Eventual consistency and node discovery

The nature of the Distributed Hash Table makes it eventually consistent. Nodes are given an initial peer list from their bootstrap node. As more nodes join the network they discover them through store and query messages as, if they do not have the requesting node in their peer list, they add these nodes to their peer list. Additionally, to supplant this, a periodic re-discovery process runs where nodes query their original bootstrap node for any peers that joined after them. This introduces extra computation onto bootstrap nodes but this can be balanced by carefully choosing more computationally powerful devices in the network to be bootstrap nodes.

4 Implementation

The following section details the implementation of the project. It includes the structure of the codebase, database schema and key algorithms. Following this it highlights key challenges and optimisations that were made during the development of the project.

4.1 Tools and technologies

TDS was implemented primarily in Go (module `tds`, Go 1.24), chosen for its lightweight concurrency model, strong standard library support for networking, and straightforward deployment as static binaries. The core system components (server, client proxy, DHT logic, registry, transport handlers, and support utilities) are all written in Go and built from the same codebase.

The development stack combines both in-memory and persistent data paths. For high-performance lookup, the server uses in-memory structures for active task mappings and round-robin state. For persistence and operational continuity, PostgreSQL is used through the Go SQL driver (github.com/lib/pq). This hybrid approach supports fast hot-path query handling while retaining a durable backing store for wider task spaces and restart scenarios.

At the interaction layer, the project supports UDP and TCP transports, with optional TLS/mutual TLS configuration for authenticated and encrypted communication in centralised deployments. Message exchange between services and the local proxy uses JSON for interoperability and easier debugging during integration testing.

For operator visibility, the server includes a terminal dashboard implemented with `tview/tcell`. This provides a lightweight monitoring interface without requiring a browser-based control plane and was useful during iterative performance and behaviour tuning.

To support realistic deployment and verification, the project includes container and orchestration assets (Dockerfiles and Kubernetes manifests), cross-platform automation scripts (PowerShell and shell), and dedicated integration test harnesses under `test_scripts/` and `test_environment/`. Together, these tools enabled repeatable local testing, multi-process simulation, and environment-specific validation for both centralised and P2P operating modes.

4.2 Codebase structure

The codebase is organised to separate executable entrypoints, reusable packages, deployment assets, and system-level test environments. This allowed me to develop centralised and P2P features in parallel while keeping the service-facing interface stable.

At the top level, `cmd/` contains runnable programs. The most important binaries are `cmd/server` for centralised operation and `cmd/client-proxy` for the local proxy process used by services in both ar-

chitectures. Additional command folders such as `cmd/demo`, `cmd/client_demo`, `cmd/test_client`, and `cmd/test_json_client` were used for integration validation and protocol probing during development.

Core implementation logic is located in `pkg/`, with package boundaries aligned to architectural concerns:

- `pkg/registry`: task-to-service registration, lookup, heartbeat handling, and load-balancing state.
- `pkg/store`: persistence and cache support for the centralised server path.
- `pkg/transport`: UDP/TCP communication handlers and message routing.
- `pkg/client`: service-facing helper functions used by applications to call REGISTER and QUERY.
- `pkg/dht`: peer management, key placement, and DHT operations for P2P mode.
- `pkg/firewall`: routing-rule filtering used by firewall-mode in centralised deployments.

Infrastructure and reproducibility assets are versioned alongside the source. Kubernetes manifests are grouped in `k8s/`, container definitions are under `cmd/*/Dockerfile`, and supporting scripts are in `scripts/`. End-to-end scenario testing is concentrated in `test_scripts/` and `test_environment/`, while the realistic transport simulation used in evaluation is provided in `Simulation/`. This structure made it possible to validate behaviour from unit-level package logic through to multi-process deployment scenarios using the same codebase.

4.3 Database schema

The centralised mode persists registry state in PostgreSQL using a single primary table, `services`. The design deliberately keeps the relational model compact: one row represents one service instance for one task, keyed by (`task`, `address`). This supports fast upsert-style registration and straightforward query aggregation while avoiding joins on the lookup path.

The effective schema used by runtime migrations is shown below.

```
CREATE TABLE IF NOT EXISTS services (  
  id SERIAL PRIMARY KEY,  
  task TEXT NOT NULL,  
  address TEXT NOT NULL,  
  last_heartbeat TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(),  
  query_count BIGINT NOT NULL DEFAULT 0,  
  capacity INTEGER NOT NULL DEFAULT 1,  
  is_active BOOLEAN NOT NULL DEFAULT TRUE,  
  created_at TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(),  
  updated_at TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(),  
  UNIQUE(task, address)  
);
```

This schema is accompanied by indexes on `task`, `last_heartbeat`, and `is_active`. Together these support the three dominant database operations in the system: task-based lookups, heartbeat timeout cleanup, and filtering of active versus inactive rows.

Two implementation decisions are important for system behaviour:

- **Idempotent registration via upsert.** Register operations use `INSERT ... ON CONFLICT (task, address) DO UPDATE` so heartbeats and capacity updates refresh existing rows without duplicate records.
- **Soft cleanup.** Expired services are marked `is_active = FALSE` rather than deleted. This preserves historical rows for observability and audit-style analysis while keeping active-query paths efficient.

The `query_count` and `capacity` columns support weighted load balancing and usage analytics. In practice, this allows the server to combine performance-oriented in-memory selection with periodic or direct persistence of usage counters, maintaining both runtime speed and durable operational state.

4.4 Functional testing

Comprehensive functional testing was conducted across both centralised and P2P modes, validating core operations including service registration, heartbeat handling, query routing, load balancing, and cleanup semantics. Test coverage includes unit tests for individual packages, integration tests for multi-component scenarios, and property-based testing for protocol correctness. Detailed test results and harnesses are provided in `test_scripts/` and `test_environment/` directories within the repository.

4.5 Integration testing: Realistic application scenario

To validate practical utility, a high-fidelity simulation of a train ticketing system (inspired by the operational problem that motivated this work) was implemented. This simulation includes station computers, payment gates, fare calculation services, and multiple ticket distribution endpoints, all interconnected through TDS service discovery. The system demonstrated seamless communication across components, with the only troubleshooting required being initial server startup—validating the robustness of both the registry logic and transport mechanisms. The complete system architecture and component interactions are documented in the High-Level Design document available in the repository.

5 Technical highlights and optimisations

Here i will lay out key insights into areas of my project that presented key challenges to me. It details how I overcame them and, in particular, where I decided to cease further work in the scope of the project timeframe.

5.1 Cache and concurrency performance

The hybrid cache-plus-database design created an unexpected performance bottleneck. During performance testing, I discovered that cache *misses* were actually faster than *hits*. This counterintuitive result revealed a concurrency pathology: while queries should be dominated by reads, every query also incremented round-robin state and query counters—both write operations requiring mutual exclusion. Even in the cache path (highest performance case), readers were competing for a write lock held during query count and round-robin updates, causing severe contention.

The solution involved two complementary optimisations:

5.1.1 Batched query count updates

Query counts were used only to determine which tasks should be cached, not for client-facing operations. This decoupling made it safe to batch counter updates asynchronously. Instead of writing to the database on every query, counters are accumulated in-memory and flushed to the database only before cache warming intervals. This eliminated one of the two write operations from the query hot path.

5.1.2 Lock-free round-robin with atomics

The round-robin load-balancing index still required update on every query to track which service should be selected next for each task. Replacing the mutex-protected counter with atomic integers allowed lock-free, wait-free updates to this state. Atomic operations in Go (via `atomic.Int64`) use CPU-level compare-and-swap instructions and do not require OS-level mutual exclusion. For a high-frequency operation like updating the round-robin index on thousands of queries per second, this eliminates the scheduler overhead of lock acquisition, contention detection, and context switching. The result is that incrementing the atomic counter is orders of magnitude faster than acquiring a mutex, performing the update, and releasing the lock.

With both optimisations in place, the cache became measurably faster than the database backend, validating the intuition that cache hits should outperform cache misses. This exercise demonstrated how

subtle concurrency bottlenecks can emerge even in simple designs and highlighted the importance of profiling to identify where lock-free atomics and asynchronous batching are most valuable.

5.1.3 Peer-to-Peer Challenges

Having implemented and tested the P2P mode for this project, I found that it struggled to scale with large numbers of nodes. The issue pertained to the fact that, as the network grew, patterns of individual nodes communicating emerged. Many nodes were rarely contacted and therefore their routing tables were not updated with the latest information. In these cases the centralised approach was far more effective. What the network gained in robustness against node failure, it lost with scale.

5.1.4 Firewall-Mode

Firewall mode was a feature I implemented with the goal of allowing the TDS system (in centralised mode) to only serve up services that were routable for the requestor. This proved to be a significant challenge as the only method I devised was to import firewall rule file from the system administrator of any given application. This was deemed unacceptable as the separation of the TDS' routing table from the source of truth (the firewall) was a recipe for inconsistency. The feature was therefore not developed further but included as an artifact as further work could be done to improve the implementation.

6 Conclusion and Discussion

6.1 Future Work

Several enhancements would improve TDS capabilities:

Replication: Add multi-master replication for centralised mode high availability. Raft consensus could provide strong consistency while tolerating failures.

Health checking: Implement active health probes instead of passive heartbeats. Server could periodically verify service reachability and automatically deregister failed instances.

Service metadata: Support arbitrary key-value tags for services enabling filtering queries (e.g., "version=2.0", "data-centre=us-east").

Multi-data-centre: Extend P2P mode with geographic awareness for locality-based routing.

Metrics export: Integrate with Prometheus for comprehensive monitoring and alerting.

REST API: Add HTTP REST interface alongside binary protocol for broader language support and web-based tooling.

6.2 Conclusion

This project has achieved the goals outline in project objectives. It provides a versatile achitecture that provides a tradeoff between correctness and decentralisation. The client proxy is a stable interface to all services running on a device enabling complete transparency to the task discovery architecture in place. With well established query, registry and reponse formats it can be easily integrated to any service. the Client package provides functions that can replace any previously existing hard coded endpoint with a call.

Centralised mode offers the benefits of a dedicated device for task discovery. In instances with generous amounts of memory, it can be run entirely from memory without disadvantage as all data must be kept fresh through heartbeats so application restarts cause minimal disruption. In conjunction with a database the server retains functionality with greater heartbeat timeouts and broad task space with a high performance cache that keeps highly queried data for fast access. The performance of these operating modes have been tested and the benefits are noticeable. The firewall functionality ensures that responses to queries are usable by a service according to a set of provided routing rules. As discussed

in the implementation section, this project does not aim to replace network security management so the routing is a best effort mechanism from a set of imported rules. A further key note is that, in hybrid mode, since both the firewall rules and tasks are stored in databases, multiple instances of this server can be present on a network to reduce load on a single server but it introduces some inconsistency with cached tasks data. Since all register operations are sent directly to the database, the impact is minimal on the correct listing of tasks.

Peer to peer mode offers fault tolerance and decentralisation at the cost of consistency and additional computation on client devices. The removal of a central authority ensures there is no single point of failure for the network and the K-Nearest Neighbours algorithm provides protection against both peer failure and inconsistency in peer knowledge for operations. As discussed in the implementation section, while this architecture is excellent for smaller deployments where it would be inefficient to have a dedicated device for service discovery, it does not scale very effectively due to the gossip protocol for peer discovery. This is an acceptable issue however as the multimodal nature of this project gives system integrators the flexibility to choose a configuration that suits their needs.

The load balancing algorithm employed in both architectures has been discussed in depth. Once again it provides flexibility for services running on both powerful and lightweight machines to receive proportional traffic. In combination with firewall routing in centralised mode it becomes a product that can be integrated into existing industry systems seamlessly. For example while building the simulation environment, I had a plethora of difficulties designing the complex system itself however the only area which presented no issues was the distribution of data for processing across the system. The project slotted into the system perfectly.

In summation, the task discovery system has gone beyond my expectations when I chose this project in objection to the component that inspired it. It has become a product that serves its purpose so well I have struggled to stop developing it further. The list of future work scratches the surface of the further potential this project possesses, there are many algorithmic and architectural improvements that can be made still but I am proud of the features and quality of the product that I have been consistently developing throughout the year.

References

- [1] Stoica, I., Morris, R., Karger, D., Kaashoek, M. F., & Balakrishnan, H. (2001). *Chord: A scalable peer-to-peer lookup service for internet applications*. ACM SIGCOMM Computer Communication Review, 31(4), 149-160.
- [2] HashiCorp. (2021). *Consul: Service Mesh for Microservices*. <https://www.consul.io/>
- [3] CoreOS. (2020). *etcd: A distributed, reliable key-value store*. <https://etcd.io/>
- [4] Hunt, P., Konar, M., Junqueira, F. P., & Reed, B. (2010). *ZooKeeper: Wait-free coordination for internet-scale systems*. USENIX Annual Technical Conference, 10(8), 9.
- [5] Karger, D., Lehman, E., Leighton, T., Panigrahy, R., Levine, M., & Lewin, D. (1997). *Consistent hashing and random trees: Distributed caching protocols for relieving hot spots on the World Wide Web*. ACM Symposium on Theory of Computing, 654-663.
- [6] Newman, S. (2015). *Building Microservices: Designing Fine-Grained Systems*. O'Reilly Media.
- [7] Donovan, A. A., & Kernighan, B. W. (2015). *The Go Programming Language*. Addison-Wesley Professional.
- [8] Cardellini, V., Colajanni, M., & Yu, P. S. (1999). *Dynamic load balancing on web-server systems*. IEEE Internet Computing, 3(3), 28-39.
- [9] Ratnasamy, S., Francis, P., Handley, M., Karp, R., & Shenker, S. (2001). *A scalable content-addressable network*. ACM SIGCOMM Computer Communication Review, 31(4), 161-172.
- [10] Maymounkov, P., & Mazieres, D. (2002). *Kademlia: A peer-to-peer information system based on the XOR metric*. International Workshop on Peer-to-Peer Systems (IPTPS), 53-65.
- [11] DeCandia, G., Hastorun, D., Jampani, M., Kakulapati, G., Lakshman, A., Pilchin, A., Sivasubramanian, S., Voshall, P., & Vogels, W. (2007). *Dynamo: Amazon's highly available key-value store*. ACM SIGOPS Operating Systems Review, 41(6), 205-220.
- [12] Das, A., Gupta, I., & Motivala, A. (2002). *SWIM: Scalable weakly-consistent infection-style process group membership protocol*. IEEE/IFIP International Conference on Dependable Systems and Networks, 303-312.
- [13] Ongaro, D., & Ousterhout, J. (2014). *In search of an understandable consensus algorithm (extended version)*. USENIX Annual Technical Conference, 305-319.
- [14] Cheshire, S., & Krochmal, M. (2013). *Multicast DNS* (RFC 6762). IETF.
- [15] Cheshire, S., & Krochmal, M. (2013). *DNS-Based Service Discovery* (RFC 6763). IETF.
- [16] Kubernetes Authors. (2026). *Kubernetes Services, Load Balancing, and Networking: Service*. <https://kubernetes.io/docs/concepts/services-networking/service/>
- [17] Kubernetes Authors. (2024). *Kubernetes Network Policies*. <https://kubernetes.io/docs/concepts/services-networking/network-policies/>
- [18] Istio Authors. (2025). *Traffic Management Concepts*. <https://istio.io/latest/docs/concepts/traffic-management/>
- [19] HashiCorp. (2026). *Consul Discover Services Overview*. <https://developer.hashicorp.com/consul/docs/discover>
- [20] etcd Authors. (2025). *etcd API guarantees*. https://etcd.io/docs/v3.6/learning/api_guarantees/
- [21] Al-Shaer, E., & Hamed, H. (2009). *Firewall policy queries*. IEEE Transactions on Parallel and Distributed Systems, 20(6), 766-777.
- [22] Jeffrey, A., & Samak, T. (2009). *Model checking firewall policy configurations*. IEEE International Symposium on Policies for Distributed Systems and Networks, 60-67.