Task Discovery Service (TDS): A Dual-Architecture Service Registry for Distributed Systems

Final Year Project Report

February 2026

1 Introduction and Motivation

In large distributed systems, one weak point can force the whole platform to slow down: static service routing. The core issue in the system I worked on at Cubic Transportation Systems was not input handling, but rigid, hard-coded output endpoints. Although the component accepted requests from a dynamic number of devices, every routing change required redevelopment and full regression testing. This made scaling expensive and slow, and over time the original resolvers had to be replaced with load balancers as demand outgrew the fixed design. This project began from that operational bottleneck and asks how service discovery can be made both dynamic and scalable.

1.1 Problem Statement

The problem lay in two key aspects

1. The component was not dynamic even though it performed the same function for any given reuqest from the devices beneath it.
2. The server was mainly overloaded becasue it accepted the data from the requesting device and passed it on to the resolver. The data was never operated upon or even observed by the component.

I proposed to address these issues by designing an architecture that would perform and scale with the system.

1.2 Research Objectives

This project addresses the service discovery problem by developing the Task Discovery Service (TDS), a flexible registry system that supports both centralised and peer-to-peer (P2P) operational modes within a unified architecture. The primary objectives are:

1. Design and implement a dual-mode service registry supporting both centralised and P2P architectures
2. Develop efficient algorithms for service registration, query routing, and load balancing
3. Create a comprehensive monitoring interface for real-time system visualisation
4. Evaluate performance characteristics under various network conditions and load patterns
5. Provide multiple transport and security options including UDP, TCP, and mutual TLS authentication

1.3 Contribution

The main contribution of this work is a production-ready service discovery system that eliminates the operational mode dichotomy present in existing solutions. TDS allows development teams to begin with a simple centralised deployment and seamlessly transition to distributed P2P operation as scale requirements increase, without changing client code or protocol definitions. This architectural flexibility, combined with comprehensive security features and real-time monitoring capabilities, represents a significant advancement in practical service discovery systems.

2 Project Background and Context

2.1 Service Discovery in Distributed Systems

Service discovery allows one part of a distributed system to find another at runtime, avoiding hard-coded addresses in an environment where instances are constantly restarted, scaled, or replaced. Most solutions fall into three groups.

2.1.1 Centralised registry approaches

Centralised approaches maintain a single authoritative catalogue. Systems such as etcd use Raft consensus to provide strongly consistent control-plane state [3,13,20]. Kubernetes builds on this by keeping stable service identities while EndpointSlices track changing pods [16]; Consul follows a similar pattern with added health checks [2,19]. The main strength is clarity and ease of observation; the main weakness is that the registry itself becomes a dependency whose failure disrupts discovery.

2.1.2 Decentralised approaches

Decentralised approaches distribute lookup data across peers. Chord and Kademlia showed that hashing keys to node IDs allows efficient $O(\log n)$ routing with no central authority [1,10]; Dynamo extended this towards high availability under network instability [11]. Membership and failure detection are handled by gossip protocols such as SWIM [12], while mDNS and DNS-SD serve smaller local networks without any dedicated server [14, 15]. The benefit is resilience to single-node failure; the cost is weaker global consistency as peer knowledge can diverge.

2.1.3 Traffic-management and policy-enforcement approaches

A third family sits above discovery and focuses on choosing between instances and controlling which connections are permitted. Istio applies load-balancing policies (round-robin, weighted routing, circuit breaking) through sidecar proxies [18], while Kubernetes NetworkPolicy restricts ingress and egress at the pod level [17]. These tools offer fine-grained control but introduce additional infrastructure and spread discovery, routing, and policy across multiple layers.

2.2 Security, Reachability, and Policy Correctness

Discovery is only useful if the returned endpoints are actually reachable. Work on firewall policy analysis—detecting rule anomalies [21] and verifying intended behaviour [22]—reinforces that a discovery system must account for the active security policy, not just record which services exist.

2.3 How TDS Differs, and Its Strengths and Weaknesses

Existing solutions require an early commitment to one architecture. TDS instead supports both centralised and peer-to-peer operation behind a single client-facing proxy interface: services issue the same

REGISTER and QUERY messages regardless of which backend is active, so the deployment topology can change without altering application code.

This reduces architectural lock-in and lowers adoption cost, while still allowing each mode to play to its strengths—centralised for simplicity and governance, P2P for resilience and removal of a central dependency. The trade-off is added implementation complexity and the fact that neither underlying architecture's limitations disappear: the centralised mode still depends on a registry server, and the P2P mode is naturally weaker on global consistency. TDS exposes both and gives the operator a choice rather than hiding those trade-offs.

3 Design

3.1 services, tasks, and the client proxy

The diagram below gives an overview of a client machine in this system. With multiple services registering and querying services to the client proxy which is the local endpoint of this solution.

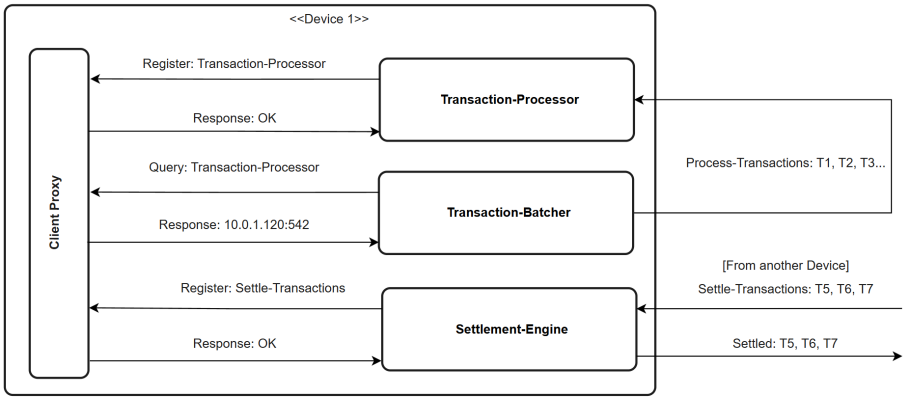

Figure 1: Client machine service registration and discovery via client proxy

In this architecture, services running on a machine will query their local instance of the client proxy. As far as the services are concerned, that is all that needs to be done for them to advertise their service along with sending periodic heartbeats. The client proxy will also handle any requests these services have for where to send their data once they have finished the scope of their operations on it. The client proxy is the interface behind which all the mechanisms of this project lie.

3.1.1 Message format

All communication between local services (clients) and the client proxy uses JSON messages in both centralised and P2P modes. The two shared commands are `REGISTER` and `QUERY`.

JSON requests (service → client proxy)

```
{"cmd":"REGISTER","task":"payment","address":"10.0.1.5:8080","capacity":10}
{"cmd":"QUERY","task":"payment"}
```

JSON response format

```
{"status":"OK"} //Register - OK
{"status":"OK","address":"10.0.1.5:8080"} //Successful query
{"status":"NOTFOUND"} // queried task not found
{"status":"ERR","error":"task required"} //example error: no task in query
    request
```

These message types are intentionally minimal so the same client-side behaviour works regardless of whether the proxy is currently forwarding to a central server or to the P2P registry.

3.2 Centralised mode

We now move outside the scope of the single machine registering and querying services to a distributed system containing many such machines. The below figure demonstrates the role the server plays in centralised mode.

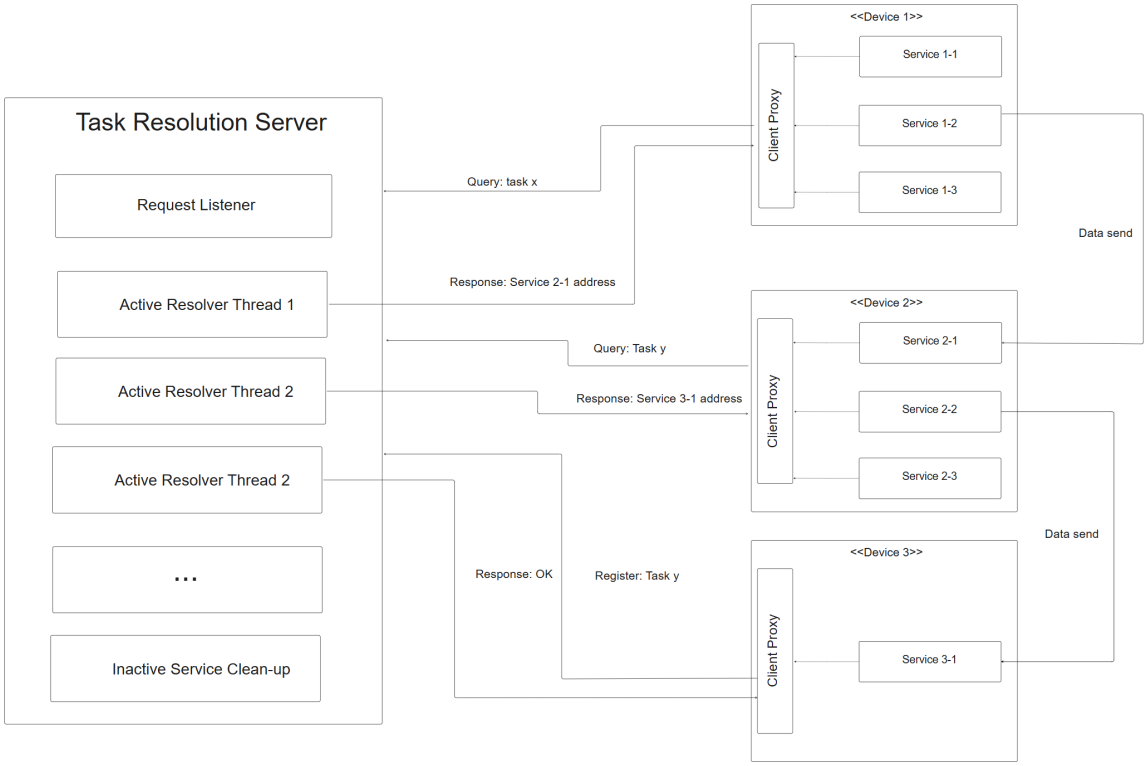

Figure 2: Client machines operating in server architecture

3.2.1 load balancing

The server load balances services using a round-robin mechanism for all services under a task. This ensures even distribution of traffic to each service. Configurable on top of this is a weighted round robin mechanism using each registered service's capacity as a weighting. This ensures that services running on less powerful machines are not swamped by being routed the same volume of traffic as a more powerful variant. During the heartbeat message, a service advertises its current capacity to inform load balancing.

3.2.2 Hybrid storage: cache and persistence

The centralised server uses a hybrid storage strategy that balances performance with operational durability. At the core is an in-memory cache that holds a subset of frequently queried tasks. Upon a cache miss or at regularised intervals, the server warms the cache by querying a persistent PostgreSQL backend. This approach achieves low-latency query response for hot tasks while maintaining durable, databaseable state for recovery and operational continuity.

In practice, the server maintains a configurable cache size and uses query frequency to determine which tasks remain cached. After query count thresholds or time-based intervals, the cache is refreshed by querying the database for the current service entries. This ensures that infrequently accessed tasks do

not consume valuable in-memory space, and that heartbeat updates in the database are reflected in the next cache warm before they are served to clients. The hybrid approach proved essential to achieving good query latency without sacrificing the ability to scale the deployment by adding a database backend and persisting across restart events.

3.2.3 Firewall-mode

In many real use cases of this project, while all machines are able to query and register with the TDS, not all services are allowed to communicate with each other. This is enforced by networking rules that are created and managed by a central authority in the system. The purpose of firewall-mode is to cross-reference a query for a service with the list of ip-address:ports that the requestor can communicate with. This way the server will only return a service address that the client can use.

3.3 Peer-to-Peer mode

The diagram below gives an overview of the peer to peer architecture. It shows how tasks are registered and then stored on the correct node in the network. It also demonstrates service queries and then the separate sending of data outside of the architecture for the process. In contrast to the centralised mode,

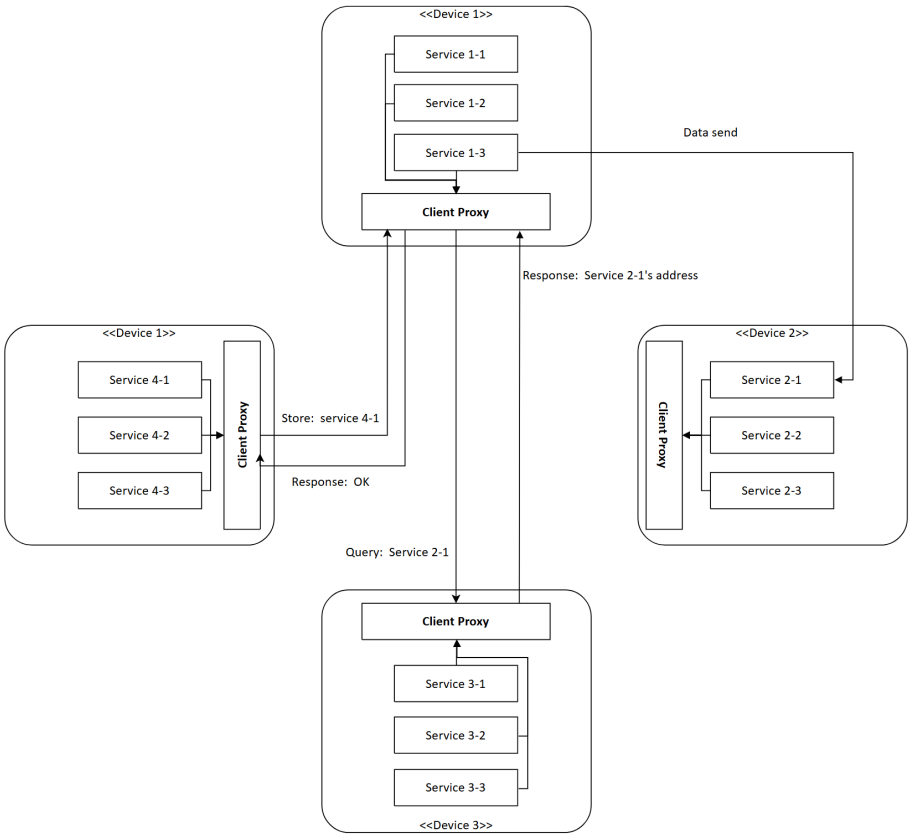

Figure 3: Client machines operating in P2P architecture

all functionality of this architecture runs from the client-proxy. The P2P network is formed from one "bootstrap" node. All subsequent nodes that join are given a pre-existing node in the network as a bootstrap. When a node joins, it hashses its ip address to receive their node ID. This ID is capped at 256 bits forming a ring network.

3.3.1 Distributed Hash Table (DHT)

The DHT is used to store the task:address mappings in this architecture. When a service registers a task with their client-proxy, the proxy hashes the service name. It then searches its list of peers to find the node ID that is closest to this hash. Once that node is found the client-proxy sends a STORE request to it with the service details. The node that receives this store message compares its peer list to the task hash to ensure it is the most appropriate candidate. If so then it stores the task in memory similarly to the cache in centralised mode.

A query follows the same path: the client-proxy hashes the query task, sends a query request to the node with the closest ID. If that node has that service stored it returns an address, if not then it returns NOTFOUND.

3.3.2 Eventual consistency and node discovery

The nature of the Distributed Hash Table makes it eventually consistent. Nodes are given an initial peer list from their bootstrap node. As more nodes join the network they discover them through store and query messages as, if they do not have the requesting node in their peer list, they add these nodes to their peer list. Additionally, to suppliment this, a perodic re-discovery process runs where nodes query their original bootstrap node for any peers that joined after them. This introduces extra computation onto bootstrap nodes but this can be balanced by carefully choosing more computationally powerful devices in the network to be bootstrap nodes.

4 Implementation

The following section details the implementation of the project. It includes the structure of the codebase, database schema and key algorithms. Following this it highlights key challenges and optimisations that were made during the development of the project.

4.1 Tools and technologies

TDS was implemented primarily in Go (module `tds`, Go 1.24), chosen for its lightweight concurrency model, strong standard library support for networking, and straightforward deployment as static binaries. The core system components (server, client proxy, DHT logic, registry, transport handlers, and support utilities) are all written in Go and built from the same codebase.

The development stack combines both in-memory and persistent data paths. For high-performance lookup, the server uses in-memory structures for active task mappings and round-robin state. For persistence and operational continuity, PostgreSQL is used through the Go SQL driver (`github.com/lib/pq`). This hybrid approach supports fast hot-path query handling while retaining a durable backing store for wider task spaces and restart scenarios.

At the interaction layer, the project supports UDP and TCP transports, with optional TLS/mutual TLS configuration for authenticated and encrypted communication in centralised deployments. Message exchange between services and the local proxy uses JSON for interoperability and easier debugging during integration testing.

For operator visibility, the server includes a terminal dashboard implemented with `tview/tcell`. This provides a lightweight monitoring interface without requiring a browser-based control plane and was useful during iterative performance and behaviour tuning.

To support realistic deployment and verification, the project includes container and orchestration assets (Dockerfiles and Kubernetes manifests), cross-platform automation scripts (PowerShell and shell), and dedicated integration test harnesses under `test_scripts/` and `test_environment/`. Together, these tools enabled repeatable local testing, multi-process simulation, and environment-specific validation for both centralised and P2P operating modes.

4.2 Codebase structure

The codebase is organised to separate executable entrypoints, reusable packages, deployment assets, and system-level test environments. This allowed me to develop centralised and P2P features in parallel while keeping the service-facing interface stable.

At the top level, `cmd/` contains runnable programs. The most important binaries are `cmd/server` for centralised operation and `cmd/client_proxy` for the local proxy process used by services in both architectures. Additional command folders such as `cmd/demo`, `cmd/client_demo`, `cmd/test_client`, and `cmd/test_json_client` were used for integration validation and protocol probing during development.

Core implementation logic is located in `pkg/`, with package boundaries aligned to architectural concerns:

- `pkg/registry`: task-to-service registration, lookup, heartbeat handling, and load-balancing state.
- `pkg/store`: persistence and cache support for the centralised server path.
- `pkg/transport`: UDP/TCP communication handlers and message routing.
- `pkg/client`: service-facing helper functions used by applications to call REGISTER and QUERY.
- `pkg/dht`: peer management, key placement, and DHT operations for P2P mode.
- `pkg/firewall`: routing-rule filtering used by firewall-mode in centralised deployments.

Infrastructure and reproducibility assets are versioned alongside the source. Kubernetes manifests are grouped in `k8s/`, container definitions are under `cmd/*/Dockerfile`, and supporting scripts are in `scripts/`. End-to-end scenario testing is concentrated in `test_scripts/` and `test_environment/`, while the realistic transport simulation used in evaluation is provided in `Simulation/`. This structure made it possible to validate behaviour from unit-level package logic through to multi-process deployment scenarios using the same codebase.

4.3 Database schema

The centralised mode persists registry state in PostgreSQL using a single primary table, `services`. The design deliberately keeps the relational model compact: one row represents one service instance for one task, keyed by `(task, address)`. This supports fast upsert-style registration and straightforward query aggregation while avoiding joins on the lookup path.

The effective schema used by runtime migrations is shown below.

```
CREATE TABLE IF NOT EXISTS services (
    id SERIAL PRIMARY KEY,
    task TEXT NOT NULL,
    address TEXT NOT NULL,
    last_heartbeat TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(),
    query_count BIGINT NOT NULL DEFAULT 0,
    capacity INTEGER NOT NULL DEFAULT 1,
    is_active BOOLEAN NOT NULL DEFAULT TRUE,
    created_at TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(),
    updated_at TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(),
    UNIQUE(task, address)
);
```

This schema is accompanied by indexes on `task`, `last_heartbeat`, and `is_active`. Together these support the three dominant database operations in the system: task-based lookups, heartbeat timeout cleanup, and filtering of active versus inactive rows.

Two implementation decisions are important for system behaviour:

- **Idempotent registration via upsert.** Register operations use `INSERT ... ON CONFLICT (task, address) DO UPDATE` so heartbeats and capacity updates refresh existing rows without duplicate records.

- **Soft cleanup.** Expired services are marked `is_active = FALSE` rather than deleted. This preserves historical rows for observability and audit-style analysis while keeping active-query paths efficient.

The `query_count` and `capacity` columns support weighted load balancing and usage analytics. In practice, this allows the server to combine performance-oriented in-memory selection with periodic or direct persistence of usage counters, maintaining both runtime speed and durable operational state.

4.4 Functional testing

Comprehensive functional testing was conducted across both centralised and P2P modes, validating core operations including service registration, heartbeat handling, query routing, load balancing, and cleanup semantics. Test coverage includes unit tests for individual packages, integration tests for multi-component scenarios, and property-based testing for protocol correctness. Detailed test results and harnesses are provided in `test_scripts/` and `test_environment/` directories within the repository.

4.5 Integration testing: Realistic application scenario

To validate practical utility, a high-fidelity simulation of a train ticketing system (inspired by the operational problem that motivated this work) was implemented. This simulation includes station computers, payment gates, fare calculation services, and multiple ticket distribution endpoints, all interconnected through TDS service discovery. The system demonstrated seamless communication across components, with the only troubleshooting required being initial server startup—validating the robustness of both the registry logic and transport mechanisms. The complete system architecture and component interactions are documented in the High-Level Design document available in the repository.

5 Technical highlights and optimisations

Here i will lay out key insights into areas of my project that presented key challenges to me. It details how I overcame them and, in particular, where I decided to cease further work in the scope of the project timeframe.

5.1 Cache and concurrency performance

The hybrid cache-plus-database design created an unexpected performance bottleneck. During performance testing, I discovered that cache *misses* were actually faster than *hits*. This counterintuitive result revealed a concurrency pathology: while queries should be dominated by reads, every query also incremented round-robin state and query counters—both write operations requiring mutual exclusion. Even in the cache path (highest performance case), readers were competing for a write lock held during query count and round-robin updates, causing severe contention.

The solution involved two complementary optimisations:

5.1.1 Batched query count updates

Query counts were used only to determine which tasks should be cached, not for client-facing operations. This decoupling made it safe to batch counter updates asynchronously. Instead of writing to the database on every query, counters are accumulated in-memory and flushed to the database only before cache warming intervals. This eliminated one of the two write operations from the query hot path.

5.1.2 Lock-free round-robin with atomics

The round-robin load-balancing index still required update on every query to track which service should be selected next for each task. Replacing the mutex-protected counter with atomic integers allowed

lock-free, wait-free updates to this state. Atomic operations in Go (via `atomic.Int64`) use CPU-level compare-and-swap instructions and do not require OS-level mutual exclusion. For a high-frequency operation like updating the round-robin index on thousands of queries per second, this eliminates the scheduler overhead of lock acquisition, contention detection, and context switching. The result is that incrementing the atomic counter is orders of magnitude faster than acquiring a mutex, performing the update, and releasing the lock.

With both optimisations in place, the cache became measurably faster than the database backend, validating the intuition that cache hits should outperform cache misses. This exercise demonstrated how subtle concurrency bottlenecks can emerge even in simple designs and highlighted the importance of profiling to identify where lock-free atomics and asynchronous batching are most valuable.

5.1.3 Peer-to-Peer Challenges

During testing, a routing inconsistency became observable in practice as the network grew. The DHT initially assigned each task to a single responsible node using `FindClosestNode`, which performs a binary search on a sorted ring of hashed node IDs. The problem emerged because nodes build their peer lists through gossip: a node only learns about a peer when it is directly contacted. In larger networks, less-trafficked nodes were rarely contacted and their peer lists fell behind. As a result, two nodes with divergent peer views would disagree on which node was closest to a given task hash. A service could register on node A (which considered itself responsible), while a query arriving at node B would compute a different "closest" node and either forward the query incorrectly or return an empty result, even though the data existed in the network. This inconsistency was visible during multi-node integration tests: queries that should have resolved would intermittently return nothing, and the failure rate increased with network size.

To address this, single-node responsibility was replaced with *k-nearest neighbour replication*. The `FindKClosestNodes` function computes the k nodes with the smallest ring distance to a task hash and registrations are written to all k of them. Lookups first attempt local storage, then check whether the querying node is among the k-nearest; if it is not, it forwards the query to all k responsible nodes in turn and returns the first non-empty response (`forwardLookup`). Because the data is replicated across k nodes, a query only fails if all k copies are simultaneously absent or unreachable. This made the system tolerant of the routing-table divergence that caused the original inconsistency: even if one node's peer list is stale and it targets the wrong primary, one of the remaining $k - 1$ replicas will respond correctly. The trade-off is k-fold write amplification, but for a service-discovery workload this is acceptable because registrations are infrequent compared to queries, and heartbeats keep the replicated entries fresh. What the network gained in robustness against node failure and stale peer knowledge, it lost somewhat in write overhead; at large enough scales the centralised approach remains more predictable, which is why TDS keeps both modes available.

5.1.4 Firewall-Mode

Firewall mode was a feature I implemented with the goal of allowing the TDS system (in centralised mode) to only serve up services that were routable for the requestor. This proved to be a significant challenge as the only method I devised was to import firewall rule file from the system administrator of any given application. This was deemed unacceptable as the separation of the TDS' routing table from the source of truth (the firewall) was a recipe for inconsistency. The feature was therefore not developed further but included as an artifact as further work could be done to improve the implementation.

6 Conclusion and Discussion

6.1 Future Work

Several enhancements would improve TDS capabilities:

Replication: Add multi-master replication for centralised mode high availability. Raft consensus could provide strong consistency while tolerating failures.

Health checking: Implement active health probes instead of passive heartbeats. Server could periodically verify service reachability and automatically deregister failed instances.

Service metadata: Support arbitrary key-value tags for services enabling filtering queries (e.g., "version=2.0", "data-centre=us-east").

Multi-data-centre: Extend P2P mode with geographic awareness for locality-based routing.

Metrics export: Integrate with Prometheus for comprehensive monitoring and alerting.

REST API: Add HTTP REST interface alongside binary protocol for broader language support and web-based tooling.

6.2 Conclusion

This project has met the objectives set out at the start. TDS provides a versatile architecture that balances correctness against decentralisation, and the client proxy acts as a stable, uniform interface to all services on a device. Because the query, registration, and response formats are well-defined, any service can integrate TDS by replacing a hard-coded endpoint with a single client library call.

Centralised mode is well-suited to deployments with a dedicated discovery server. When memory is plentiful the server can operate entirely in-memory: because all records must be kept fresh via heartbeats, a restart causes minimal disruption. Pairing the server with a database extends its capacity and allows longer heartbeat timeouts, while a high-performance cache keeps the most-queried entries fast. The firewall feature adds another layer of utility by filtering query responses against imported routing rules, ensuring that only reachable endpoints are returned to a caller. As noted in the implementation discussion, TDS does not aim to replace dedicated network security tooling; the routing is best-effort. In hybrid (database-backed) mode, multiple server instances can share load because both tasks and firewall rules live in the database rather than in local memory. Cache divergence between instances is possible, but because all registrations go directly to the database the effect on task listings is small.

Peer-to-peer mode trades consistency and some per-client computation for fault tolerance and the removal of any single point of failure. The K-Nearest Neighbours algorithm makes lookups resilient to both peer failure and stale peer knowledge. As discussed in the implementation section, this mode suits smaller deployments where a dedicated discovery server would be wasteful, though its gossip-based peer discovery limits scalability at larger scales. That limitation is manageable precisely because TDS is multimodal: integrators can choose the topology that fits their environment.

The weighted round-robin load balancer used in both modes deserves a mention. It distributes traffic proportionally, so lightweight devices are not overwhelmed alongside more powerful ones. Combined with firewall-aware routing in centralised mode, it makes TDS straightforward to slot into an existing system. This was evident during the development of the simulation environment: designing the distributed system as a whole was challenging, yet routing data between its components for processing required no special effort at all.

In summary, TDS has exceeded my initial expectations. It has grown into a product that performs its intended role so effectively that I have found it hard to stop adding to it. The future work listed above only scratches the surface of what is possible; there remain meaningful algorithmic and architectural improvements still to make. I am proud of what has been built and of the consistency with which it has been developed throughout the year.

References

[1] Stoica, I., Morris, R., Karger, D., Kaashoek, M. F., & Balakrishnan, H. (2001). *Chord: A scalable peer-to-peer lookup service for internet applications*. ACM SIGCOMM Computer Communication Review, 31(4), 149-160.

[2] HashiCorp. (2021). *Consul: Service Mesh for Microservices*. https://www.consul.io/

[3] CoreOS. (2020). *etcd: A distributed, reliable key-value store*. https://etcd.io/

[4] Hunt, P., Konar, M., Junqueira, F. P., & Reed, B. (2010). *ZooKeeper: Wait-free coordination for internet-scale systems*. USENIX Annual Technical Conference, 10(8), 9.

[5] Karger, D., Lehman, E., Leighton, T., Panigrahy, R., Levine, M., & Lewin, D. (1997). *Consistent hashing and random trees: Distributed caching protocols for relieving hot spots on the World Wide Web*. ACM Symposium on Theory of Computing, 654-663.

[6] Newman, S. (2015). *Building Microservices: Designing Fine-Grained Systems*. O'Reilly Media.

[7] Donovan, A. A., & Kernighan, B. W. (2015). *The Go Programming Language*. Addison-Wesley Professional.

[8] Cardellini, V., Colajanni, M., & Yu, P. S. (1999). *Dynamic load balancing on web-server systems*. IEEE Internet Computing, 3(3), 28-39.

[9] Ratnasamy, S., Francis, P., Handley, M., Karp, R., & Shenker, S. (2001). *A scalable content-addressable network*. ACM SIGCOMM Computer Communication Review, 31(4), 161-172.

[10] Maymounkov, P., & Mazieres, D. (2002). *Kademlia: A peer-to-peer information system based on the XOR metric*. International Workshop on Peer-to-Peer Systems (IPTPS), 53-65.

[11] DeCandia, G., Hastorun, D., Jampani, M., Kakulapati, G., Lakshman, A., Pilchin, A., Sivasubramanian, S., Vosshall, P., & Vogels, W. (2007). *Dynamo: Amazon's highly available key-value store*. ACM SIGOPS Operating Systems Review, 41(6), 205-220.

[12] Das, A., Gupta, I., & Motivala, A. (2002). *SWIM: Scalable weakly-consistent infection-style process group membership protocol*. IEEE/IFIP International Conference on Dependable Systems and Networks, 303-312.

[13] Ongaro, D., & Ousterhout, J. (2014). *In search of an understandable consensus algorithm (extended version)*. USENIX Annual Technical Conference, 305-319.

[14] Cheshire, S., & Krochmal, M. (2013). *Multicast DNS* (RFC 6762). IETF.

[15] Cheshire, S., & Krochmal, M. (2013). *DNS-Based Service Discovery* (RFC 6763). IETF.

[16] Kubernetes Authors. (2026). *Kubernetes Services, Load Balancing, and Networking: Service*. https://kubernetes.io/docs/concepts/services-networking/service/

[17] Kubernetes Authors. (2024). *Kubernetes Network Policies*. https://kubernetes.io/docs/concepts/services-networking/network-policies/

[18] Istio Authors. (2025). *Traffic Management Concepts*. https://istio.io/latest/docs/concepts/traffic-management/

[19] HashiCorp. (2026). *Consul Discover Services Overview*. https://developer.hashicorp.com/consul/docs/discover

[20] etcd Authors. (2025). *etcd API guarantees*. `https://etcd.io/docs/v3.6/learning/api_guarantees/`

[21] Al-Shaer, E., & Hamed, H. (2009). *Firewall policy queries*. IEEE Transactions on Parallel and Distributed Systems, 20(6), 766-777.

[22] Jeffrey, A., & Samak, T. (2009). *Model checking firewall policy configurations*. IEEE International Symposium on Policies for Distributed Systems and Networks, 60-67.